

# Data & Things

(Spring 26)

---

Monday February 16

## Lecture 8: Classification II

Jens Ulrik Hansen

# Outline of this lecture

---

- Hyper-parameter tuning and cross-validation
- Decision trees for classification (and regression)
- Ensemble models using bagging and boosting
- Dealing with unbalanced classes
- Exercises

# Hyper-parameter tuning and cross-validation

---

- **Hyper-parameters**
  - Many machine learning *models are parametric*, in the sense of having parameters which values *we optimize through training* – for instance, the a and b of the linear regression model  $y = a + b*x$ .
  - Furthermore, many machine learning models also have *hyper-parameters*, that is, *parameters we set in advance* for each model, which are *never changed during training* of a model – for instance, the K in KNN classifiers.
  - Each setting of hyper-parameter values give rise to a different model. Finding the setting of hyper-parameters that result in the best model, is called *hyper-parameter tuning*.
  - Visually looking for the optimal K in KNN classification, was a way of doing simple hyper-parameter tuning – once we have multiple hyper-parameters however, we need a more systematic and precise way of doing hyper-parameter tuning!
- More broadly, we sometimes also want to do model selection on other things than hyper-parameters, such as variable selection


# Hyper-parameter tuning and cross-validation

---

- **Our train-test split is often not enough!**
  - We want train multiple models and pick the best one, or iteratively adjust “hyper-parameters” of our models
    - if we use the performance on the test set to compare different models, we might *risk overfitting to the test set*
    - Moreover, the performance measure on the test set is no longer a good estimate of how the model will perform on new data
  - If we have smaller datasets, or use many variables in our model, we are left with a dilemma
    - We need a large portion of the data for training, to be able to get a good model
    - We need a large portion of the data for testing, to get a good estimate of how we will perform on new unseen data
  - Finally, train-test split can be quite sensitive to the actual random split

# Hyper-parameter tuning and cross-validation


- Recall our KNN example on the diabetes example from last time...
  - We wanted to use the test accuracy to pick the right K
  - We had only 768 data points in total, so we got few data points both for training and testing no matter how we did the split
  - Different train-test split give different optimal Ks (as we shall see)
  - Can we somehow utilize all the data for training and still get a good (unbiased) estimate of how well we will perform on future unseen data?



# Hyper-parameter tuning and cross-validation


- **Train-Validation-Test split**

- Use an additional **validation set** to evaluate and compare models
  - We leave the test set as is and split the training set further into a new training set and a validation set.
  - We then train our models on the new training set
  - We evaluate, compare and chose the final model based on performance on the validation set
  - Finally, we calculate an unbiased estimate of our model performance on new unseen data by calculating the performance on the test set
  - There is no one right proportion of splits, but an often used split is: 60% - training, 20% - validation, 20% - test



# Hyper-parameter tuning and cross-validation

- ***k*-fold cross-validation** ( $k = 5$  or  $k = 10$  are popular choices)
  - Split the (train and validation) data  $D$  into  $k$  mutually exclusive subsets, each of approximately equal size:  $D_1 \dots D_k$ . Each  $D_i$  is called a *fold*.
  - Do model construction and evaluation  $k$  times. Use the *average* accuracy.
    - At the  $i$ -th iteration, use fold  $D_i$  as the validation set and the others as the training set.
- Example of standard 3-fold cross validation



# Hyper-parameter tuning and cross-validation

- **Stratified cross-validation**

- folds are stratified so that *class distribution* in each fold is approximately the same as that in the initial given data.



- **Leave-one-out cross-validation**


- $k$  is set to the number of data points/rows in the dataset
- Use it only for small sized data; otherwise, too many models to construct!
- Often 5- or 10-fold cross validation is just as good or even better

# Hyper-parameter tuning and cross-validation

- Standard 3-fold cross validation



- Stratified 3-fold cross validation




# Hyper-parameter tuning and cross-validation

---

- Cross-validation is not a way to construct an applicable model.
  - The function `cross_val_score(.)` builds multiple models *internally*, but these models are not returned.
- The purpose of cross-validation is to evaluate how well a *type* of model will generalize when it is trained on a specific dataset.
  - Model type: decision tree, random forest, KNN, SVM, ...
- Moreover, by using cross-validation, we can decide what type of model to use, and tune hyper-parameters for constructing a model
  - **Hyper-parameters**: algorithm parameters that can be set by the user before training a model.
    - E.g., `K` for KNN, `test_size` and `random_state` for `train_test_split(.)`, `gini` or `entropy` for a DT, ...
  - In contrast, **model parameters** are learned internally from training data
    - E.g., `a` and `b` coefficients in simple linear regression or logistic regression, weights in a neural network, ...

# Hyper-parameter tuning and cross-validation

- *After having used cross-validation to decide model type and tune hyper-parameters, we train the final model one last time on all the (training and validation) data (before evaluating it on the untouched test data)*



# Hyper-parameter tuning and cross-validation

---

- **Examples**

- Let us look at the notebook “Hyper-parameter tuning and cross-validation.ipynb”.

# Outline of this lecture

---


- Hyper-parameter tuning and cross-validation
- Decision trees for classification (and regression)
- Ensemble models using bagging and boosting
- Dealing with unbalanced classes
- Exercises

# Decision trees for classification



- **Decision Trees**

- This classification model is organized as a tree for decision making – it is thus called a **decision tree**.
- Internal nodes are associated with an *attribute/column* and arcs with *values* for that attribute.
- A leaf node tells the predicted class label (where the branches ends).

- Each person has attributes/columns:
  - (age, student [yes/no], credit rating)
  - p1(18, yes, fair)
  - p2(55, no, excellent)
- **Two classes**
  - **Buy computer**
  - Not buy computer



# Decision trees for classification



- **Training a decision tree model**

- Recursive binary splits top-down.
- In a 2D feature dataset, each split corresponds to drawing a horizontal or vertical line.
- Splits are chosen to minimize either *Gini index* or *Entropy* (both measures of node impurity).
- Each final region corresponds to a leaf in the decision tree.
- For each region, the predicted class corresponds to the most prevalent class – class probabilities can be obtained by noting the fraction of each class.

<https://github.com/martian1231/decisionTreeFromScratch>


# Decision trees for classification

- **Pros and Cons of decision trees**

- Pros: Easy to interpret (even by non experts) and visualize
- Pros: Does not require much preprocessing such as scaling or creation of dummy variables for categorical variables
- Pros: Highly flexible
- Cons: Can grow big and then becomes hard to interpret
- Cons: Does not have the same level of predictive accuracy as other models
- Cons: Have high variance

- **Decision trees for regression**

- The prediction is made by taking the average of values in a region
- Instead of minimizing Gini index or entropy, we minimize a local version of RSS.



# Decision trees for classification

---

- **Examples**

- Let us look at the notebook “Decision trees and ensemble models.ipynb”.

# Outline of this lecture

---

- Hyper-parameter tuning and cross-validation
- Decision trees for classification (and regression)
- Ensemble models using bagging and boosting
- Dealing with unbalanced classes
- Exercises

# Ensemble models using bagging and boosting

---

- ***Ensemble models*** combine many simple models into one single potential powerful model
- Decision trees for classification (or regression) turns out to be good candidates for such simple models
- Popular ensemble methods:
  - ***Bagging***: averaging the prediction over a set of independent classifiers
  - ***Random Forest***: a random set of decision trees
  - ***Boosting***: weighted vote with a set of dependent classifiers

# Ensemble models using bagging and boosting

---

- **Bagging: Bootstrap aggregating**

- Analogy: Diagnosis based on the *majority vote* of multiple doctors trained from samples of the same superset.
- Reduce variance of models with high variance, such as decision trees
- Training: Given a set  $D$  of  $d$  tuples, generate  $m$  new training data sets  $D_i$ 's
  - Each training set  $D_i$  of  $d'$  tuples is sampled with *replacement* from  $D$  (i.e., **bootstrap**)
  - Set  $D_i$  is used to train a classifier model  $M_i$
- Classification: classify an unknown sample  $x$  using all  $m$  models ( $M_1$  to  $M_m$ )
  - Each classifier  $M_i$  returns its class prediction
  - The bagged classifier  $M^*$  counts the votes and assigns the class with the most votes to  $x$
- Accuracy
  - Often significantly better than a single classifier derived from  $D$
  - For noisy data: not considerably worse, more robust

# Ensemble models using bagging and boosting

---

- **Random Forest (of Decision Trees)**

- An improved version of bagging
  - In addition to bootstrap sampling, we are also sampling a subset of features for each decision tree
- Formally, for each decision tree
  - We sample with replacement a subset of the original training set
  - We sample set of  $m$  of the feature variables – if there is a total of  $p$  features, a common choice for  $m$  is  $\sqrt{p}$
  - We then train the decision tree on this subset of the training dataset and only allows for split involving the sampled  $m$  features.
- The rationale behind Random Forest
  - That we also sample only a small subset of features for each decision tree makes the decision trees much more varied and uncorrelated, which in turn make the average more robust.
- Decision rule
  - To make a prediction, Random forest average over all predictions from all the trees for regression and take the most popular vote for a class among all the trees for classification

# Ensemble models using bagging and boosting

---

- **Boosting**

- Analogy: Consult several doctors, based on a combination of weighted diagnoses—weight assigned based on the *previous* diagnosis accuracy
- How boosting works?
  - Weights are assigned to each training tuple
  - A series of  $k$  classifiers is iteratively learned
  - After a classifier  $M_i$  is learned, the weights are updated to allow the subsequent classifier,  $M_{i+1}$ , to pay more attention to the training tuples that were misclassified by  $M_i$
  - The final  $M^*$  combines the votes of each individual classifier, where the weight of each classifier's vote is a function of its accuracy
- Compared with bagging: Boosting tends to have greater accuracy, but it also risks overfitting the model to misclassified data
- Popular variations of boosting: ***AdaBoost*** and ***XGBoost***

# Ensemble models using bagging and boosting

---

- **Interpretability of ensemble models**

- As we are combining many models to achieve higher accuracy, interpretability is no longer straight-forward – a classic example of the trade-off between predictive accuracy and interpretability
- However, for ensemble models based on decision trees, we can record in each decision note and for each feature variable  $X_i$  the drop in Gini index (or local RSS) and average this across all the trees in the ensemble – this will give us the ***variable importance*** of the feature variable  $X_i$ .
  - In this way, we can get a ranking of which of the features are most important relative to each other in the predictions the ensemble model makes.

# Summary on Decision trees and ensemble models

---

- Decision trees are easy to understand, easy to train, and easy to interpret
- Decision trees are flexible and easily handle both numeric and categorical variables (without scaling) and can be used for both regression and classification
- Decision trees often have limited accuracy and very high variance
- For these reasons, Decision trees are ideal as the simple models combined in ensemble models
- Random Forest is an ensemble method that often perform well and rarely overfit
- Boosting methods are more prone to overfitting, but can also achieve higher predictive accuracy than Random Forest

# Ensemble models using bagging and boosting

---

- **Examples**

- Let us look at the notebook “Decision trees and ensemble models.ipynb”.

# Outline of this lecture

---

- Hyper-parameter tuning and cross-validation
- Decision trees for classification (and regression)
- Ensemble models using bagging and boosting
- Dealing with unbalanced classes
- Exercises

# Dealing with unbalanced classes

---

- **Class imbalance**
  - Rare positive examples but numerous negative ones, e.g., medical tests, fraud detection, etc. (...or the other way around)
  - Traditional methods assume a balanced distribution of classes and equal error costs: not suitable for class-imbalanced data
  - Typical methods for imbalance data in binary class classification:
    - **Oversampling**: re-sampling of data from positive class
    - **Under-sampling**: randomly eliminate tuples from negative class
    - **Threshold-moving**: moves the decision threshold,  $t$ , so that the rare class tuples are easier to classify, and hence, less chance of costly false negative errors
    - ...
  - Still difficult for class imbalance problem on multiclass tasks

# Dealing with unbalanced classes

---

- **Class imbalance**

- Examples: Rare positive but numerous negative ones, e.g., medical tests, fraud detection, etc. (...or the other way around)
- Traditional machine learning methods assume a balanced distribution of classes and equal error costs – ***these are not suitable for class-imbalanced data!***
- Example: Imagine a classification task with a response variable that is either 1 for the subject a disease, and 0 if the subject does not have the disease.
  - If it is a rare disease, maybe only 5% of the people/data points in our training data has 1 for the response variable.
  - In this case, we can define a very simple model with 95% accuracy, if we just always predict 0!
  - The model might “think” a little along these lines, in the sense that cases with  $y=0$ , will have less effect on the training.

# Dealing with unbalanced classes

- **Typical methods to deal with unbalanced data in (binary) classification**
  - **Oversampling**: Multiple sampling with replacement of data from minority class
  - **Undersampling**: Only sample some of the data points from the majority class
  - **Combining over- and under-sampling**: Combining the two above approaches
  - **More advanced sampling technique like SMOKE:**
    - Chawla, N. V., Bowyer, K. W., Hall, L. O., & Kegelmeyer, W. P. (2002). SMOTE: synthetic minority over-sampling technique. *Journal of artificial intelligence research*, 16, 321-357. <https://www.jair.org/index.php/jair/article/view/10302>
    - **Threshold-moving**: moves the decision threshold,  $t$ , so that the rare class tuples are easier to classify, and hence, less chance of costly false negative errors
    - ...

# Dealing with unbalanced classes

---

- **Examples**

- Let us look at the notebook “Dealing with unbalanced classes.ipynb”.

# Outline of this lecture

---

- Hyper-parameter tuning and cross-validation
- Decision trees for classification (and regression)
- Ensemble models using bagging and boosting
- Dealing with unbalanced classes
- Exercises

# Exercises

---

- Do Exercise 1 and 2 in the notebook “Exercises in Classification II.ipynb”